

# notebook\_aha

April 21, 2021

## 1 Code for Aggregating Heterogeneous-Agent Models with Permanent Income Shocks

*Karl Harmenberg*

*2021-04*

In this Python notebook, I describe the details associated with the computational experiments in the paper Aggregating Heterogeneous-Agent Models with Permanent Income Shocks. The source code for the underlying code is provided in the file `code_aha.py`. The computations are implemented with just-in-time compilation through [the Numba package](#). For an introduction to Python for economics (or vice versa, an introduction to quantitative economics for Pythonistas), see [QuantEcon](#).

First we import the code from `code_aha` as well as some math and plotting capabilities from `numpy` and `matplotlib`. We also import `time` because you never want to run out of time...

```
[1]: from code_aha import *
import numpy as np
from numpy.polynomial.hermite import hermgauss
import matplotlib.pyplot as plt
import matplotlib as mpl
mpl.rcParams['figure.dpi'] = 300
import seaborn as sns
sns.set_theme()
import time
```

### 1.1 Calibration and discretization of the model.

**Discretized state space:** Next, we describe the discretized grids used in the computational experiments. The grid for cash on hand consists of 300 grid points from 0.1 to 400 and the grid for end-of-period savings consists of 300 grid points from 0 to 400. The grid is finer closer to the borrowing constraint.

The method introduced in the paper, employing the permanent-income-neutral measure, permits computing model aggregates without keeping track of the permanent-income dimension. However, for comparison purposes we will also compute model aggregates without the permanent-income-neutral measure. Therefore, I introduce an equispaced grid in log permanent income from -10 to 10 with 31 grid points.

```
[2]: convexity = 2.0

mgrid = np.linspace(0.10**(1/convexity), 400**(1/convexity), 300)**convexity
dmgrid = np.hstack([mgrid[1:]-mgrid[:-1], [mgrid[-1]-mgrid[-2]]])
bgrid = np.hstack([[0.0], (np.linspace(0.10**(1/convexity), 400**(1/convexity),
↪299))**convexity])

zgrid = np.linspace(-10, 10, 31)
```

**Discretized shock distribution:** For both the permanent income shocks  $\eta$  and the transitory income shocks  $\epsilon$ , I use Gauss-Hermite quadrature with 5 points to describe the shock distributions. The variance of the log of the transitory shock is  $0.01 \times 4$  and the variance of the log of the permanent income shock is  $0.01 \times 4/11$ , following [Carroll, Slacalek, Tokuoka, and White \(QE, 2017\)](#). The shocks processes are adjusted to ensure that  $E[\eta] = E[\epsilon] = 1$ .

```
[3]: sigma_epsilon = np.sqrt(0.01*4)
sigma_eta = np.sqrt(0.01*4/11)

Nshock = 5
val, prob = hermgauss(Nshock)
val = np.sqrt(2)*val
prob = prob/np.sqrt(np.pi)
Nepsilon = Nshock
Neta = Nshock

epsilon_prob_temp = prob
epsilon_val_temp = np.exp(val*sigma_epsilon)
epsilon_val_temp = epsilon_val_temp/np.sum(epsilon_prob_temp*epsilon_val_temp)

epsilon_prob = epsilon_prob_temp
epsilon_val = epsilon_val_temp

eta_prob_temp = prob
eta_val_temp = np.exp(val*sigma_eta)/np.sum(eta_prob_temp*np.exp(val*sigma_eta))

eta_prob = eta_prob_temp
eta_val = eta_val_temp
```

Following [Carroll, Slacalek, Tokuoka, and White \(QE, 2017\)](#), household preferences are described by a discount rate  $\beta = 0.99$  and a risk aversion  $\sigma = 1.0$ .

Finally, to maintain a stationary income distribution, households die in a perpetual-youth fashion at rate 0.00625, yielding an average working life of 40 years (again, following [Carroll, Slacalek, Tokuoka, and White \(QE, 2017\)](#)). Note that, with the permanent-income-neutral measure, it is not *necessary* to have a stationary income distribution in order to compute model aggregates, however since I aim to compare with not using the permanent-income-neutral measure, I need a non-degenerate income distribution.

We take all these parameters and save them as a `problem_parameters` object.

```
[4]: = 0.99
      = 1.0

death_prob = 0.00625

params = problem_parameters( , , mgrid, bgrid, zgrid,
                             Nepsilon, epsilon_val, epsilon_prob,
                             Neta, eta_val, eta_prob, death_prob)
```

## 1.2 Optimal consumption behavior

The household problem (on recursive form) is

$$V(M, Z) = \max_C u(C) + \beta EV(M', Z')$$

subject to:

(i) the budget constraint,

$$B + C = M,$$

(ii) a borrowing constraint,

$$B \geq 0,$$

(iii) the definition of cash on hand in the next period,  $M'$ , as savings from the current period and next-period income,

$$M' = RB + Y'$$

and (iv) the income process

$$Y' = wZ'\epsilon',$$

$$Z' = \eta'Z.$$

Finally,  $\epsilon' \sim F_\epsilon$  and  $\eta' \sim F_\eta$  are i.i.d. over time.

As is well-known, if  $u(\cdot)$  is CRRA with risk aversion  $\gamma$ , then  $V$  can be written on the form  $V(A, Z) = v(A/Z)Z^{1-\gamma}$ . Denoting  $a = A/Z$ ,  $b = B/Z$ , etc., we arrive at the normalized household problem.

### 1.2.1 Normalized household problem

$$v(m) = \max_c \frac{c^{1-\gamma}}{1-\gamma} + \beta E[(\eta')^{1-\gamma} v(m')]$$

subject to

$$m' = R \frac{m - c}{\eta'} + w\epsilon',$$

$$c \leq m,$$

and the shock distributions  $\epsilon' \sim F_\epsilon$  and  $\eta' \sim F_\eta$ .

For a given wage and interest rate, the household problem is solved with a combination of the endogenous-grid method ([Carroll 2006](#)) and a matrix implementation of Howard's improvement algorithm (see p. 18 of [Moritz Kuhn's notes](#), thanks to Pontus Rendahl for teaching me this “trick”).

Below, the optimal consumption function is computed given the wage and price of a bond (as well as an initial guess for the value function, I choose the flow utility of consuming all the cash on hand as the initial guess).

### 1.2.2 The consumption function

Below, I compute the consumption function and plot it. This is the first time the code is run in this notebook, so it will take a few extra seconds for Numba to compile the code. Next time we run the code it will be much, much faster.

```
[5]: initial_v = u(params.mgrid, params.)


wage = 1.0
Q = 0.9925
R = 1/Q

consumption_function, v = compute_optimal_consumption_function(initial_v, R, wage, params)

plt.plot(params.mgrid, consumption_function, label = 'Consumption function')
plt.plot([0, 2], [0, 2], '--', color = 'k', alpha = 0.2, label = '45 degree line')
plt.xlim([0, 10])
plt.xlabel('Cash on hand')
plt.ylabel('Consumption')
plt.legend()
plt.ylim([0, 1.5])
```

Iterations needed= 12

```
[5]: (0.0, 1.5)
```



### 1.3 Computing the stationary permanent-income-weighted distribution

The function `compute_cash_on_hand_distribution` computes the stationary permanent-income-weighted distribution using the method from the paper. First, the code creates a discretized transition matrix, then the stationary distribution is computed as the eigenvector associated with eigenvalue 1. The source code is here:

```
def compute_cash_on_hand_distribution(consumption_function, params,
                                     R, wage, which = 'Marginal'):

    if which == 'Marginal':
        weighting_scheme = 'None'
    if which == 'Permanent Income Weighted':
        weighting_scheme = 'Aggregate'

    #Write down discretized transition matrix:
    transition_matrix_aggregate \
        = create_transition_matrix(consumption_function, weighting_scheme,
                                   params, R, wage)

    #Use a scipy routine to compute the eigenvector associated with
    #eigenvalue 1, i.e., the stationary distribution.
    eigv, stationary_distribution = \
        eigs(transition_matrix_aggregate.transpose(), k = 1, sigma = 1.0)
```

```

#The eigenvector that the scipy routine returns should be
#normalized to sum to 1
stationary_distribution = stationary_distribution.real.flatten()\
                        /np.sum(stationary_distribution.real)

return stationary_distribution

```

The function `transition_matrix_aggregate` takes as input the previously computed consumption function, the parameters of the problem and the prices (the wage and the interest rate). Furthermore, it takes as an input whether one should compute permanent-income-weighted dynamics or not.

Inside `transition_matrix_aggregate`, the only difference between the two options are the following lines.

```

if weighting_scheme == 'None':
    weight = prob
if weighting_scheme == 'Aggregate':
    weight = params.eta_val[eta_i]*prob

```

### 1.3.1 Computation time using the permanent-income-weighted distribution

Below, I compute the permanent-income-weighted distribution using the above method and record the computation time.

```

[6]: start = time.time()
    piw_distribution = compute_cash_on_hand_distribution(consumption_function,
                                                         params, R, wage,
                                                         which = 'Permanent Income_
↳Weighted')
    end = time.time()
    print("Computation time:", end - start)

    start = time.time()
    marginal_distribution = compute_cash_on_hand_distribution(consumption_function,
                                                             params, R, wage,
                                                             which = 'Marginal')
    end = time.time()
    print("Computation time:", end - start)

    plt.plot(params.mgrid, piw_distribution/dmgrid,
              label = 'Permanent-income-weighted distribution')
    plt.plot(params.mgrid, marginal_distribution/dmgrid,
              label = 'Marginal distribution')
    plt.legend()
    plt.xlabel('Normalized cash on hand')
    plt.xlim([0, 40])
    plt.title('Cash-on-hand distributions')
    plt.show()

```


```


plt.plot(params.mgrid-consumption_function, piw_distribution/dmgrid,
         label = 'Permanent-income-weighted distribution')
plt.plot(params.mgrid-consumption_function, marginal_distribution/dmgrid,
         label = 'Marginal distribution')
plt.legend()
plt.xlabel('Normalized end-of-period assets')
plt.xlim([0, 40])
plt.title('Asset distribution')
plt.show()

```

Computation time: 0.01095890998840332

Computation time: 0.009007930755615234





The slight jaggedness of the distributions stems from the discreteness of the shock distribution. The normalized end-of-period asset distribution also has a point mass of households with exactly zero assets.

### 1.3.2 Computation time without using the permanent-income-weighted distribution

**Take aways from this section:** 1. With 31 grid points in permanent income, the computation of the stationary distribution is slowed down by 2 orders of magnitude. Furthermore, the discrepancy between the permanent-income-weighted distribution computed with the permanent-income-neutral measure and the permanent-income-weighted distribution computed using the grid in permanent income is noticeable. 2. With 101 grid points in permanent income, the computation of the stationary distribution is slowed down by 3 orders of magnitude (i.e., a factor 1000), but the discrepancy between the two methods is small.

Below, I compute the full stationary distribution (over both permanent income and cash on hand). The code for the method is virtually identical, but now the state space is substantially larger (300 grid points in the cash-on-hand dimension and 31 grid points in the permanent-income dimension means a 31 times larger state space, a total of 9300 grid points).

```
[7]: start = time.time()
stationary_distribution_2d = \
    compute_stationary_distribution_2d(consumption_function,
                                       params, R, wage)
end = time.time()
```




```
print("Computation time:", end - start)
```

Computation time: 0.9080018997192383

The computation time went from 0.01 seconds to almost one second, i.e., a slowdown of two orders of magnitude.

Below, I plot the stationary distribution of log permanent income.


```
[8]: plt.plot(params.zgrid, np.sum(stationary_distribution_2d, axis = 0))
plt.xlabel("log permanent income")
plt.title("Permanent income distribution, 31 grid points")
plt.savefig("permanent_income_dist_31.pdf")
plt.show()
```



The stationary distribution of log permanent income does not look too good. It should look like a two-sided exponential but something funky is going on with the left tail. There is an easy fix to this problem: Increase the state space for permanent income. The only drawback is that it becomes even more computationally expensive to compute the stationary distribution.

Next, I plot the permanent-income-weighted distribution implied by the full two-dimensional stationary distribution. Ideally, the two methods of computing the permanent-income-weighted distribution, directly and indirectly, should give the same answer. Let's see if they do!

```
[9]: zstate = np.vstack([params.zgrid for k in range(params.Nm)])
plt.plot(params.mgrid, np.sum(stationary_distribution_2d*np.exp(zstate), axis =
→1)/dmgrid,
        label = 'With 31 grid points in permanent income')
plt.plot(params.mgrid, piw_distribution/dmgrid, '--',
        label = 'With permanent-income-neutral measure',
        color = 'b')
plt.xlim([0,20])
plt.title("Comparison: Permanent-income-weighted distributions")
plt.xlabel('Cash on hand')
plt.legend()
plt.savefig('discrepancy_31.pdf')
plt.show()
```



They largely do agree, although the gap between the two is visible. How do we reduce the gap? The gap is reduced if we increase the number of grid points for permanent income, but, again, this comes at an additional computational cost.

To bring home this point, let's increase the number of grid points to 101 in the permanent income dimension and redo the computations.

```
[10]: zgrid_new = np.linspace(-10, 10, 101)
params_new = problem_parameters(, , mgrid, bgrid, zgrid_new,
                                Nepsilon, epsilon_val, epsilon_prob,
```

```

        Neta, eta_val, eta_prob, death_prob)

start = time.time()
stationary_distribution_2d_new = \
    compute_stationary_distribution_2d(consumption_function,
                                      params_new, R, wage)
end = time.time()
print("Computation time:", end - start)

```

Computation time: 14.397378921508789

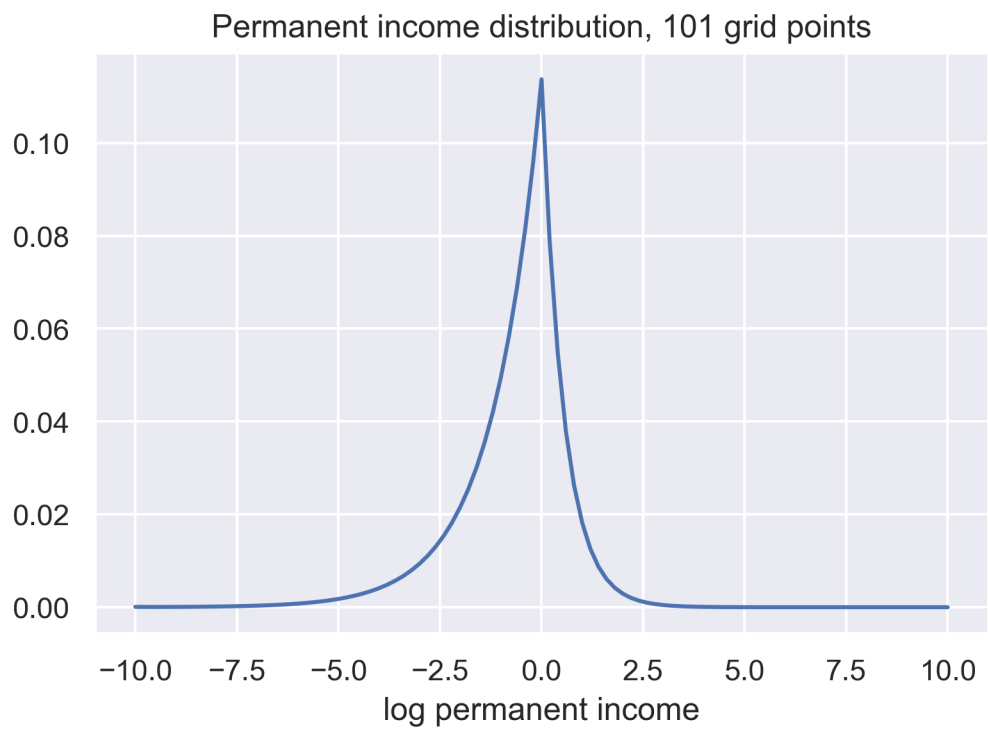
The computation time increases further, up to more than 10 seconds. Below, let's look at how the stationary distribution looks with 101 grid points in permanent income.

```

[11]: plt.plot(params_new.zgrid,
              np.sum(stationary_distribution_2d_new, axis = 0))
plt.xlabel("log permanent income")
plt.title("Permanent income distribution, 101 grid points")
plt.savefig("permanent_income_dist_101.pdf")
plt.show()

zstate_new = np.vstack([params_new.zgrid for k in range(params_new.Nm)])
plt.plot(params_new.mgrid,
         np.sum(stationary_distribution_2d_new*np.exp(zstate_new), axis = 1)/
         ↪dmgrid)
plt.plot(params.mgrid, piw_distribution/dmgrid, '--',
         color = 'b')
plt.xlim([0,20])
plt.title("Comparison: Permanent-income-weighted distributions")
plt.xlabel('Cash on hand')
plt.show()


```



The log permanent income distribution “looks good” (like a two-sided exponential distribution) and the gap between the permanent-income-weighted distribution computed using the permanent-income-neutral measure and the implied permanent-income-weighted distribution from the two-dimensional distribution is no longer visible.

To visually inspect the improvement in precision, I plot the *difference* between the two distributions, both for 31 and 101 grid points.

```
[12]: plt.plot(params.mgrid,
              np.sum(stationary_distribution_2d_new*np.exp(zstate_new), axis = 1)\
              -piw_distribution, label = 'Error with 101 grid points')
plt.plot(params.mgrid,
          np.sum(stationary_distribution_2d*np.exp(zstate), axis = 1)\
          -piw_distribution, label = 'Error with 31 grid points')
plt.xlim([0,40])
plt.legend()
plt.savefig('diff_approx.pdf')
```



With 101 grid points in the permanent income dimension, we, for all intents and purposes, recover the permanent-income-weighted distribution that we obtained using the permanent-income-neutral measure, albeit with a computational speed which is three orders of magnitude slower.

#### 1.4 Solving an Aiyagari model using the permanent-income-neutral measure

We now take our aggregation method and use it to solve an Aiyagari model. The central function is `compute_implied_savings` that takes a level of the capital stock and return the implied steady state level of savings. The equilibrium is obtained when the capital stock equals the savings.

The code for the function is below. It computes the prices implied by the capital stock, computes the consumption function, computes the permanent-income-weighted distribution and then finally backs out the savings from the permanent-income-weighted distribution.

```
def compute_implied_savings(Kguess, initial_v, params, alpha, delta):

    R = (alpha*Kguess**(alpha-1)+(1-delta))/(1-params.death_prob)
    wage = (1-alpha)*Kguess**alpha
    inner_start = time.time()
    consumption_function, v = \
        compute_optimal_consumption_function(initial_v, R, wage, params)
    inner_mid_point = time.time()
    piw_distribution = \
        compute_cash_on_hand_distribution(consumption_function, params,
                                          R, wage,
                                          which = 'Permanent Income Weighted')

    inner_end = time.time()

    b_end_of_period = params.mgrid-consumption_function
    K = np.sum(b_end_of_period*piw_distribution)
    print("Difference:", Kguess-K)
    print("K:", K)
    print("K/Y:", K**(1-alpha))
    print("Consumption computation time =", inner_mid_point - inner_start)
    print("Distribution computation time =", inner_end - inner_mid_point)
    print("*"*50)
    return K, v, piw_distribution
```

Since we are interested in computational performance, it also prints how long time it takes to compute optimal consumption and the permanent-income-weighted distribution.

I take this function and put it inside a simple while loop that implements Broyden's method.

The capital share  $\alpha$  for the production function  $Y = K^\alpha L^{1-\alpha}$  and the depreciation rate  $\delta$  are (again) taken from [Carroll, Slacalek, Tokuoka, and White \(QE, 2017\)](#).

```
[13]: start = time.time()

initial_v = u(mgrid, params.)

#Production is Cobb Douglas Y = K^*L^(1- ) with capital share .
    = 0.36
#Capital depreciates at rate
    = 0.025

Kold = 60.0
K = 50.0
v = initial_v
```

```

Kold_implied, v, piw_distribution = \
    compute_implied_savings(Kold, initial_v, params, , )
initial_v = v
K_implied, v, piw_distribution = \
    compute_implied_savings(K, initial_v, params, , )

while np.abs(K_implied-K) > 1e-10:
    slope = (np.log(K_implied)-np.log(K)-\
             np.log(Kold_implied)+np.log(Kold))/(K-Kold)
    Kguess = K - (np.log(K_implied)-np.log(K))/slope

    Kguess_implied, v, piw_distribution = \
        compute_implied_savings(Kguess, initial_v, params,
                                , )

    Kold_implied, Kold = K_implied, K
    K_implied, K = Kguess_implied, Kguess
    initial_v = v

end = time.time()

print("Total time:", end-start)

```

```

Iterations needed= 12
Difference: 41.85896146752771
K: 18.141038532472287
K/Y: 6.3906314640193
Consumption computation time = 0.04633188247680664
Distribution computation time = 0.009104251861572266
*****
Iterations needed= 8
Difference: -18.84828737124937
K: 68.84828737124937
K/Y: 15.005539366832721
Consumption computation time = 0.03266406059265137
Distribution computation time = 0.008416891098022461
*****
Iterations needed= 8
Difference: -6.158226474330462
K: 58.268198439245396
K/Y: 13.485784578278984
Consumption computation time = 0.0314481258392334
Distribution computation time = 0.008786916732788086
*****
Iterations needed= 6
Difference: 0.7501027142170713

```

```

K: 52.49197436633395
K/Y: 12.61419269957123
Consumption computation time = 0.024514198303222656
Distribution computation time = 0.008367061614990234
*****
Iterations needed= 6
Difference: -0.031475950519208595
K: 53.1459553939937
K/Y: 12.714548417622263
Consumption computation time = 0.024296283721923828
Distribution computation time = 0.009080648422241211
*****
Iterations needed= 5
Difference: -0.00016403258727848424
K: 53.11975760581366
K/Y: 12.710536852348604
Consumption computation time = 0.02322983741760254
Distribution computation time = 0.008301019668579102
*****
Iterations needed= 4
Difference: 3.5206831228151714e-08
K: 53.11962033456598
K/Y: 12.710515830661546
Consumption computation time = 0.017433881759643555
Distribution computation time = 0.00877523422241211
*****
Iterations needed= 3
Difference: -6.608047442568932e-13
K: 53.11962036402328
K/Y: 12.71051583517263
Consumption computation time = 0.013300895690917969
Distribution computation time = 0.00828695297241211
*****
Total time: 0.28582000732421875

```

The equilibrium of the Aiyagari model is found in less than one third of a second.

## 1.5 Simulating with stochastic simulation

Another way of computing the stationary distribution, with less programming and thinking overhead, but with worse computational performance is by Monte Carlo simulation.

The key code for the simulation is the following:

```

def simulate_agent_one_period(m, P, consumption_function,
                             params,
                             R, wage,
                             distorted_probabilities):

```



```

#Transitory shock
epsilon_shock = random_choice(params.epsilon_val, params.epsilon_prob)

if distorted_probabilities == False:
    #Permanent-income shock if not using the
    #permanent-income-neutral measure
    eta_shock = random_choice(params.eta_val,
                              params.eta_prob)
else:
    #If using the permanent-income-neutral measure, the shock
    #probability distribution is adjusted
    eta_shock = random_choice(params.eta_val,
                              params.eta_prob*params.eta_val)

#Death shock
death_shock = random_choice(np.array([0,1]),
                             np.array([1-params.death_prob, params.death_prob]))

if death_shock == 0:
    b = m-linint(m, params.mgrid, consumption_function)
    m_new = wage*epsilon_shock + R*b/eta_shock

    if distorted_probabilities == False:
        P_new = eta_shock*P
    else:
        P_new = 1.0 #With permanent-income-neutral measure, the permanent
                     #income should not be updated.
else:
    b = 0.0
    m_new = wage*epsilon_shock
    P_new = 1.0

return m_new, P_new, b

```

The code below takes one or two minutes to run. It simulates 100 iterations of the simulation of one agent for 1 000 000 periods.

```

[14]: start = time.time()

K_nonstochastic = K

R = ( *K_nonstochastic**(-1)+(1- ))/(1-params.death_prob)
wage = (1- )*K_nonstochastic**

consumption_function, v = compute_optimal_consumption_function(initial_v, R,
↪wage, params)

```

```

b_end_of_period = params.mgrid-consumption_function

Klist_old_method = []
distorted_probabilities = False

Ndraws = 100

for seed in range(Ndraws):
    m = 15.0
    P = 1.0
    Nperiods = int(1e6)

    mlist, Plist, blist = \
        simulate_agent_many_periods(m, P, consumption_function, Nperiods,
                                    params,
                                    R, wage,
                                    distorted_probabilities, seed)

    Klist_old_method.append(np.mean(blist*Plist))

end = time.time()
print("*"*30)
print("Computation complete")
print("Total computation time:", end-start)


plt.hist(mlist, density = True, bins = 100, weights = Plist)
plt.plot(params.mgrid, piw_distribution/dmgrid)
plt.savefig("stochastic_simulation_piw.pdf")
plt.show()

```

```

Iterations needed= 1
*****
Computation complete
Total computation time: 84.49607825279236

```



```
[15]: start = time.time()

Klist_new_method = []
distorted_probabilities = True

for seed in range(Ndraws):
    m = 15.0
    P = 1.0
    Nperiods = int(1e6)

    mlist, Plist, blist = \
        simulate_agent_many_periods(m, P, consumption_function, Nperiods,
                                    params,
                                    R, wage,
                                    distorted_probabilities, seed)


    Klist_new_method.append(np.mean(blist*Plist))

end = time.time()
print("*"*30)
print("Computation complete")
print("Total computation time:", end-start)

plt.hist(mlist, density = True, bins = 100, weights = Plist)
plt.plot(params.mgrid, piw_distribution/dmgrid)
```

```
plt.show()
```

```
*****  
Computation complete  
Total computation time: 98.0776960849762
```



## 1.6 Stochastic results for the paper

Below, I summarize the above

```
[16]: print("Comparison")  
print("W/o PIN measure:", np.mean(Klist_old_method), np.std(Klist_old_method)/  
      ↳ np.sqrt(Ndraws))  
print("W/ PIN measure:", np.mean(Klist_new_method), np.std(Klist_new_method)/np.  
      ↳ sqrt(Ndraws))  
print(np.mean(Klist_old_method)/np.mean(Klist_new_method),  
      np.std(Klist_old_method)/np.std(Klist_new_method))
```

Comparison

W/o PIN measure: 52.93853389671755 0.14961385356743498

W/ PIN measure: 52.851884405206626 0.06230294546254591

1.001639477806441 2.401392943089295

## 1.7 Non-stochastic results for the paper

Below, I compute the non-stochastic results for the paper.

```

[17]: b_end_of_period = params.mgrid-consumption_function

start = time.time()
piw_distribution = compute_cash_on_hand_distribution(consumption_function,
                                                    params, R, wage,
                                                    which = 'Permanent Income_
↳Weighted')
end = time.time()
print(""*30)
print("With PIN measure")
print("Computation time:", end - start)
print("Savings:", np.sum(b_end_of_period*piw_distribution))

print(""*30)
print("With 31 grid points")
start = time.time()
stationary_distribution_2d =_
↳compute_stationary_distribution_2d(consumption_function, \
                                     params, R, wage)
end = time.time()
print("Computation time:", end - start)

zstate = np.vstack([params.zgrid for k in range(params.Nm)])
print("Savings:",
      np.sum(b_end_of_period*np.sum(stationary_distribution_2d*np.exp(zstate),_
↳axis = 1)))

print(""*30)
print("With 101 grid points")
start = time.time()
stationary_distribution_2d_new =_
↳compute_stationary_distribution_2d(consumption_function, \
                                     params_new,_
↳R, wage)
end = time.time()
print("Computation time:", end - start)
print("Savings:",
      np.sum(b_end_of_period*np.sum(stationary_distribution_2d_new*np.
↳exp(zstate_new), axis = 1)))

```

```

*****
With PIN measure
Computation time: 0.010415792465209961
Savings: 53.11962036402264
*****
With 31 grid points

```

Computation time: 1.1111879348754883

Savings: 50.864529936550596

\*\*\*\*\*

With 101 grid points


Computation time: 11.582109212875366

Savings: 53.108711064931825

Figure for the paper:

```
[18]: piw_distribution = compute_cash_on_hand_distribution(consumption_function,
                                                         params, R, wage,
                                                         which = 'Permanent Income_
→Weighted')
marginal_distribution = compute_cash_on_hand_distribution(consumption_function,
                                                         params, R, wage,
                                                         which = 'Marginal')

plt.plot(params.mgrid-consumption_function, piw_distribution/dmgrid,
         label = 'Permanent-income-weighted distribution')
plt.plot(params.mgrid-consumption_function, marginal_distribution/dmgrid,
         label = 'Marginal distribution')
plt.legend()
plt.xlabel('Normalized end-of-period assets')
plt.xlim([0, 200])
plt.title('Asset distribution')
plt.savefig('histogram.pdf')
plt.show()
```



Finally, we compute the equilibrium consumption function:

```
[19]: plt.plot(params.mgrid, consumption_function, label = 'Consumption function')
plt.plot([0, 5], [0, 5], '--', color = 'k', alpha = 0.2, label = '45 degree_
↪line')
plt.xlim([0, 10])
plt.xlabel('Cash on hand')
plt.ylabel('Consumption')
plt.legend()
plt.ylim([0, 3])
plt.savefig('consumption_function.pdf')
```

